

Spring 2024: CS5720 Neural Networks & Deep Learning – ICP-8

Assignment – 7

Name: Mahendra Sriram Betha

700757819

Github Link: https://github.com/mahendrasrirambetha/NN_Assignment7

Video Link: https://drive.google.com/file/d/1Gzh5bhKDVJy4P17simN8s_JHnwa-SYIW/view?usp=drive_link

Use Case Description:

LeNet5, AlexNet, Vgg16, Vgg19

1. Training the model
2. Evaluating the model

Programming elements:

1. About CNN
2. Hyperparameters of CNN
3. Image classification with CNN

In class programming:

1. Tune hyperparameter and make necessary addition to the baseline model to improve validation accuracy and reduce validation loss.
2. Provide logical description of which steps lead to improved response and what was its impact on architecture behavior.
3. Create at least two more visualizations using matplotlib (Other than provided in the source file)
4. Use dataset of your own choice and implement baseline models provided.
5. Apply modified architecture to your own selected dataset and train it.
6. Evaluate your model on testing set.
7. Save the improved model and use it for prediction on testing data
8. Provide plot of confusion matrix
9. Provide Training and testing Loss and accuracy plots in one plot using subplot command and history object.
10. Provide at least two more visualizations reflecting your solution.
11. Provide logical description of which steps lead to improved response for new dataset when compared with baseline model and enhance architecture and what was its impact on architecture behavior.

Chrome File Edit View History Bookmarks Profiles Tab Window Help

localhost:8889/notebooks/icp8_700757819.ipynb

zoom M Paused Finish update

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions.

Jupyter icp8_700757819 Last Checkpoint: 26 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help

In [1]:

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
import tensorflow_datasets as tfds
from tensorflow.keras.datasets import cifar100

from tensorflow.keras.optimizers import RMSprop
from keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Dropout, BatchNormalization
%matplotlib inline
```

Extract data and train and test dataset

In [2]:

```
cifar100 = tf.keras.datasets.cifar100
(X_train, Y_train), (X_test, Y_test) = cifar100.load_data()
```

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>
169801437/169801437 [=====] - 6s 0us/step

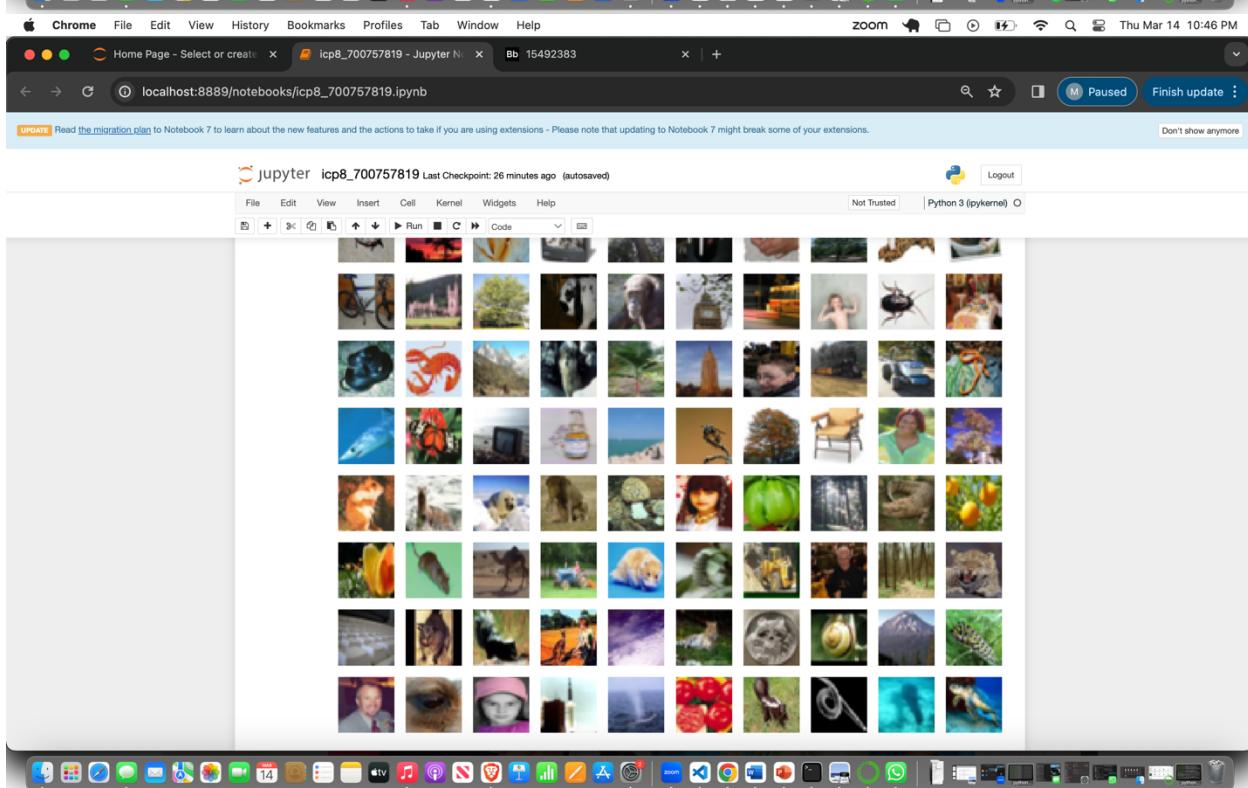
In [27]:

```
classes = ['apple', 'aquarium_fish', 'babby', 'bear', 'beaver', 'bed', 'bee', 'beetle', 'bicycle', 'bottle',
```

Let's look into the dataset images

In [3]:

```
plt.figure(figsize = (16,16))
for i in range(100):
    plt.subplot(10,10,i+1)
    plt.axis('off')
    plt.imshow(X_train[i], cmap = 'gray')
```



Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:47 PM

localhost:8889/notebooks/icp8_700757819.ipynb

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions.

Don't show anymore

Jupyter icp8_700757819 Last Checkpoint: 26 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) Logout

Training , Validating and Splitting trained and tested data

```
In [4]: from sklearn.model_selection import train_test_split
x_train, x_val, y_train, y_val = train_test_split(X_train,Y_train,test_size=0.2)

In [6]: from keras.utils import to_categorical
y_train = to_categorical(y_train, num_classes = 100)
y_val = to_categorical(y_val, num_classes = 100)

In [7]: print(x_train.shape)
print(y_train.shape)
print(x_val.shape)
print(y_val.shape)
print(X_test.shape)
print(Y_test.shape)

(40000, 32, 32, 3)
(40000, 100)
(10000, 32, 32, 3)
(10000, 100)
(10000, 32, 32, 3)
(10000, 1)

In [8]: train_datagen = ImageDataGenerator(
    preprocessing_function = tf.keras.applications.vgg19.preprocess_input,
    rotation_range=10,
    zoom_range = 0.1,
    width_shift_range = 0.1,
    height_shift_range = 0.1,
    shear_range = 0.1,
    horizontal_flip = True
)
train_datagen.fit(x_train)

val_datagen = ImageDataGenerator(preprocessing_function = tf.keras.applications.vgg19.preprocess_input)
val_datagen.fit(x_val)

In [9]: from keras.callbacks import ReduceLROnPlateau
learning_rate_reduction = ReduceLROnPlateau(monitor='val_accuracy',
                                             patience=3,
                                             verbose=1,
```

Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:47 PM

localhost:8889/notebooks/icp8_700757819.ipynb

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions.

Don't show anymore

Jupyter icp8_700757819 Last Checkpoint: 27 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) Logout

We have used only 16 layers out of 19 layers in the CNN

```
In [10]: vgg_model = tf.keras.applications.VGG19(
    include_top=False,
    weights=None,
    input_shape=(32,32,3),
)
vgg_model.summary()
Model: "vgg19"
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[None, 32, 32, 3]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160
block4_conv2 (Conv2D)	(None, 4, 4, 512)	23590800

Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:47 PM

localhost:8889/notebooks/icp8_700757819.ipynb Bb 15492383

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions. Don't show anymore

Jupyter icp8_700757819 Last Checkpoint: 27 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) Logout

```

block4_pool (MaxPooling2D) (None, 2, 2, 512) 0
block5_conv1 (Conv2D) (None, 2, 2, 512) 2359808
block5_conv2 (Conv2D) (None, 2, 2, 512) 2359808
block5_conv3 (Conv2D) (None, 2, 2, 512) 2359808
block5_conv4 (Conv2D) (None, 2, 2, 512) 2359808
block5_pool (MaxPooling2D) (None, 1, 1, 512) 0
=====
Total params: 20,024,384
Trainable params: 20,024,384
Non-trainable params: 0

In [14]: model = tf.keras.Sequential()
model.add(vgg19)
model.add(Flatten())
model.add(Dense(1024, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(1024, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dense(256, activation = 'relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))
model.add(Dense(100, activation = 'softmax'))

model.summary()
Model: "sequential_1"
Layer (type)          Output Shape         Param #
=====
```

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten_1 (Flatten)	(None, 512)	0
dense_4 (Dense)	(None, 1024)	525312

Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:47 PM

localhost:8889/notebooks/icp8_700757819.ipynb Bb 15492383

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions. Don't show anymore

Jupyter icp8_700757819 Last Checkpoint: 27 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) Logout

```

batch_normalization_3 (BatchNormalization) 4096
dense_5 (Dense) (None, 1024) 1049600
batch_normalization_4 (BatchNormalization) 4096
dense_6 (Dense) (None, 256) 262400
batch_normalization_5 (BatchNormalization) 1024
dropout_1 (Dropout) (None, 256) 0
dense_7 (Dense) (None, 100) 25700
=====
Total params: 21,896,612
Trainable params: 21,892,004
Non-trainable params: 4,608

In [15]: optimizer = tf.keras.optimizers.SGD(learning_rate = 0.001, momentum = 0.9)
model.compile(optimizer=optimizer,
              loss='categorical_crossentropy',
              metrics=['accuracy'])

In [16]: history = model.fit(
            train_datagen.flow(x_train, y_train, batch_size = 128),
            validation_data = val_datagen.flow(x_val,y_val, batch_size = 128),
            epochs = 10,
            verbose = 1,
            callbacks = [learning_rate_reduction]
        )
Epoch 1/10
313/313 [=====] - 44s 124ms/step - loss: 5.2395 - accuracy: 0.0140 - val_loss: 4.5463 - val_accuracy: 0.0205 - lr: 0.0010
Epoch 2/10
315/315 [=====] - 38s 121ms/step - loss: 5.0465 - accuracy: 0.0181 - val_loss: 4.5047 - va
```

Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:47 PM

localhost:8889/notebooks/cp8_700757819.ipynb Bb 15492383

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions.

Don't show anymore

Jupyter ipc8_700757819 Last Checkpoint: 27 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) O

In [17]:

```
Epoch 1/10
313/313 [=====] - 44s 124ms/step - loss: 5.2395 - accuracy: 0.0140 - val_loss: 4.5463 - va
l_accuracy: 0.0205 - lr: 0.0010
Epoch 2/10
313/313 [=====] - 38s 121ms/step - loss: 5.0465 - accuracy: 0.0181 - val_loss: 4.5047 - va
l_accuracy: 0.0240 - lr: 0.0010
Epoch 3/10
313/313 [=====] - 36s 116ms/step - loss: 4.9621 - accuracy: 0.0179 - val_loss: 4.7844 - va
l_accuracy: 0.0117 - lr: 0.0010
Epoch 4/10
313/313 [=====] - 36s 115ms/step - loss: 4.8896 - accuracy: 0.0191 - val_loss: 4.5794 - va
l_accuracy: 0.0206 - lr: 0.0010
Epoch 5/10
313/313 [=====] - 36s 116ms/step - loss: 4.7757 - accuracy: 0.0202 - val_loss: 4.4016 - va
l_accuracy: 0.0302 - lr: 0.0010
Epoch 6/10
313/313 [=====] - 36s 116ms/step - loss: 4.6797 - accuracy: 0.0253 - val_loss: 4.3797 - va
l_accuracy: 0.0325 - lr: 0.0010
Epoch 7/10
313/313 [=====] - 35s 113ms/step - loss: 4.6021 - accuracy: 0.0245 - val_loss: 4.3531 - va
l_accuracy: 0.0295 - lr: 0.0010
Epoch 8/10
313/313 [=====] - 36s 115ms/step - loss: 4.5665 - accuracy: 0.0258 - val_loss: 4.3573 - va
l_accuracy: 0.0308 - lr: 0.0010
Epoch 9/10
313/313 [=====] - 36s 116ms/step - loss: 4.5199 - accuracy: 0.0262 - val_loss: 4.3059 - va
l_accuracy: 0.0348 - lr: 0.0010
Epoch 10/10
313/313 [=====] - 36s 113ms/step - loss: 4.4793 - accuracy: 0.0281 - val_loss: 4.3472 - va
l_accuracy: 0.0322 - lr: 0.0010
```

Out[17]: <matplotlib.legend.Legend at 0x7f82e645a190>

Chrome File Edit View History Bookmarks Profiles Tab Window Help zoom 🔍 ⚡ Thu Mar 14 10:48 PM

localhost:8889/notebooks/cp8_700757819.ipynb Bb 15492383

UPDATE Read the migration plan to Notebook 7 to learn about the new features and the actions to take if you are using extensions - Please note that updating to Notebook 7 might break some of your extensions.

Don't show anymore

jupyter ipc8_700757819 Last Checkpoint: 27 minutes ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python 3 (ipykernel) O

Out[17]:

In [18]:

```
acc = history.history['accuracy']
val_acc = history.history['val_accuracy']

plt.figure()
plt.plot(acc,color = 'purple',label = 'Training Accuracy')
plt.plot(val_acc,color = 'blue',label = 'Validation Accuracy')
plt.legend()
```

Out[18]: <matplotlib.legend.Legend at 0x7f82e61d22e0>

Out[19]:

In [19]:

```
X_test = tf.keras.applications.vgg19.preprocess_input(X_test)
y_pred = np.argmax(model.predict(X_test), axis=-1)
y_pred[10]

313/313 [=====] - 3s 9ms/step
```

Out[19]: array([51, 77, 96, 5, 87, 77, 77, 77, 52], dtype=int32)

In [20]:

```
from sklearn.metrics import confusion_matrix, accuracy_score
print("Testing Accuracy : ", accuracy_score(y_test, y_pred))
```

Testing Accuracy : 0.8283

Chrome File Edit View History Bookmarks Profiles Tab Window Help

localhost:8889/notebooks/icp8_700757819.ipynb

In [19]: `X_test = tf.keras.applications.vgg19.preprocess_input(X_test)`
`y_pred = np.argmax(model.predict(X_test), axis=-1)`
`y_pred[:10]`
313/313 [=====] - 3s 9ms/step
Out[19]: array([53, 77, 36, 5, 87, 77, 77, 52, 82])

In [20]: `from sklearn.metrics import confusion_matrix, accuracy_score`
`print('Testing Accuracy : ', accuracy_score(y_test, y_pred))`
Testing Accuracy : 0.8383

In [21]: `cm = confusion_matrix(y_test, y_pred)`
Out[21]: array([[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]])

In [30]: `import itertools`
`def plot_confusion_matrix(cm, classes,`
 `normalize=True,`
 `title='Confusion matrix',`
 `cmap=plt.cm.Greens:`
 `'''`
 `This function prints and plots the confusion matrix.`
 `Normalization can be applied by setting 'normalize=True'.`
 `plt.imshow(cm, interpolation='nearest', cmap=cmap)`
 `plt.title(title)`
 `plt.colorbar()`
 `tick_marks = np.arange(len(classes))`
 `plt.xticks(tick_marks, classes, rotation=30)`
 `plt.yticks(tick_marks, classes)`
 `if normalize:`
 `cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]`
 `print('Normalized confusion matrix')`
 `else:`
 `print('Confusion matrix, without normalization')`
 `# print(cm)`
 `thresh = cm.max() / 2.`
 `for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):`
 `plt.text(j, i, cm[i, j],`
 `horizontalalignment="center",`
 `color="white" if cm[i, j] > thresh else "black")`
 `plt.tight_layout()`
 `plt.ylabel('True label')`
 `plt.xlabel('Predicted label')`

